

The Sotheby's Luxury Week Playbook

*Cross-category performance, premium drivers, and the HK–NY buyer corridor
2021–2026*

BY ZHUOYAO DENG

Market Overview

Sotheby's Luxury Week has become the firm's highest-velocity sale calendar, compressing watches, jewels, handbags, wine, and modern collectibles into a single concentrated selling window. Across 1,180 lots sold between 2021 and 2026, this study models what drives a lot to clear above estimate, and finds the answer is strikingly consistent across otherwise unrelated categories.

01 / MARKET

Watches leads on premium without the volatility.

Across the study window Watches held a median hammer-to-mid-estimate ratio of 1.41x against a sell-through rate near 95%. Modern Collectibles reached higher in its best seasons but also fell further in its worst. Watches rarely had a weak quarter, which is usually what you see when there are more buyers than good material to go around.

02 / CATEGORY

The ceiling is set by rarity, not by name.


In every category model the largest premium attaches to scarcity or documented history, not to the maker's name. A Himalaya crocodile Birkin runs about 56% above its ordinary counterpart, an exotic Paul Newman Daytona more than 200%, a Karuizawa bottling close to 60%. The name and the age are priced in before bidding even starts. What moves the hammer is whatever is left that can't be reproduced.

03 / CORRIDOR

Where a lot sells depends on what it is.

Comparing 80 lots that sold in both Hong Kong and New York shows there is no single better venue, only a better venue for a given thing. Hermès handbags clear around 8% higher in Hong Kong, while vintage Rolex Daytona clears closer to 17% higher in New York. Tested together the two effects nearly cancel, which is why the routing question only becomes useful once the categories are pulled apart.

I. Macro




Source: Sotheby's Past Auctions (sothebys.com); author's analysis. Figures are illustrative of the analytical method.

The starting question for any sale calendar is which categories are carrying the room. The heatmap tracks median premium ratio, hammer over mid-estimate, for each anchor category across every Luxury Week season from 2021 to 2026. The more telling pattern is not which season ran hot but how differently the categories behave from one another.

Watches holds a tight band of strong results season after season. Modern Collectibles posts the highest individual cells in the grid yet also the coldest, swinging between blockbuster and ordinary depending on what crossed the block that quarter. A strong Modern Collectibles season usually comes down to one or two exceptional consignments, not broad demand, and it pays to read it that way before drawing conclusions.

Premium and sell-through do not always move together. A category can clear nearly every lot while only modestly beating estimate, which tends to mean estimates were set conservatively against steady demand. Watches manages the rarer combination, holding sell-through high while still leading on premium, the profile of a category where buyers outnumber the good material coming to market.

II. Hermès Birkin & Kelly




Source: Sotheby's Past Auctions (sothebys.com); author's analysis. Figures are illustrative of the analytical method.

Working from 295 sold lots, the model regresses log-premium on the five characteristics a catalogue records for every bag: leather, hardware, size, condition, and model. The reference point is the bag one sees most often at auction, a 35cm Birkin in Togo leather with gold hardware in Excellent condition, so each coefficient reads as the premium a feature adds against that ordinary baseline.

Two features pull away from the rest. A Himalaya Niloticus crocodile finish carries a premium near 56% over the same bag in Togo, and diamond-set hardware sits close behind around 45%. The two are hard to separate statistically, but both stand well clear of everything beneath them, so the honest reading is that they share the top tier rather than that one outranks the other.

Condition is where the more useful lesson sits, because its effect is not symmetric. A bag graded Good rather than Excellent gives back close to a third of its premium, more than Pristine adds at the other end. The downside of a condition slip is bigger than the upside. For estimate-setting that means pricing anything below Excellent conservatively, less to maximise a single hammer than to protect sell-through and the momentum of the sale.

III. HK–NY Buyer Corridor



Note: Each point is a matched pair: the same reference or model sold in both HK and NY within a 30% mid-estimate band. Points above the dashed line indicate HK paid more; points below indicate NY paid more.

Source: Sotheby's Past Auctions (sothebys.com); author's analysis.


Source: Illustrative matched-pair sample; author's analysis. Buyer-share figures cited below: Sotheby's press release, December 2022.

There is a structural reason a corridor exists. Asian clients account for roughly 40% of all spend across Sotheby's core luxury auctions, the highest share of any region, and their weight is uneven across categories: about 34% of handbag buyers and a comparable share in watches. That buyer geography suggests, though it does not by itself prove, that a category whose collector base sits disproportionately in Asia should find deeper competition in Hong Kong, while one whose deepest bench is American should hold its strongest bids in New York.

To test it, the analysis pairs comparable lots, the same reference or model in similar condition and within a 30% band of one another's mid-estimate, and compares the premium each realised in Hong Kong against New York. Points above the parity line did better in Hong Kong, points below did better in New York. The categories separate by direction, and the direction is the one the buyer geography predicted: handbags around 8% higher in Hong Kong, vintage Daytona closer to 17% stronger in New York, Patek near parity.

A pooled test across all pairs returns a p-value near 0.10, which a strict reading would call inconclusive. The reason is not weak data but offsetting signals that cancel when averaged into one number. The pooled figure is the wrong thing to read; the split underneath it is the point, and it lines up with where each category's buyers actually are.

IV.1 Vintage Rolex Daytona




Source: Sotheby's Past Auctions (sothebys.com); author's analysis. Figures are illustrative of the analytical method.

The vintage Daytona market is organised almost entirely around the dial. An exotic Paul Newman dial lifts hammer premium by more than 200% over an otherwise identical standard-dial reference, a coefficient several times larger than anything else in the model and large enough that it effectively defines two separate markets that happen to share a case shape.

Full documentation and named provenance matter too, around 47% and 49%, but they are second-order next to the dial. For a watch sale the catalogue order is obvious: put the Paul Newman lots first, shoot the dial properly, and let everything else sit behind them.

IV.2 Patek Philippe Calatrava



Source: Sotheby's Past Auctions (sothebys.com); author's analysis. Figures are illustrative of the analytical method.

Calatrava premium is led by two features that are difficult to rank against each other. A salmon dial and a stainless-steel case each carry something near 50% to 58%, with overlapping confidence intervals, so the data treats them as a shared top tier. Both are scarce for the same reason: Patek built the Calatrava overwhelmingly in precious metal with conventional dials, so steel and salmon are the exceptions a collector waits years to see.

The more practical lever is the next tier down. Original papers add roughly 36% and a Patek archive extract about 24%, and unlike a rare case or dial, you can actually go and get them. No one can produce a steel Calatrava on demand, but a specialist can write to Patek for an extract, which makes the paperwork the most dependable way to lift a Calatrava result.

Further categories. The same model is applied to DRC Burgundy and Japanese whisky in the full online report. In Burgundy, format and direct-from-Domaine provenance lead, the latter worth about 22% over a private-cellar source at the same vintage. In whisky, a closed distillery dominates: Karuizawa runs close to 60% over a mainstream distillery of comparable age, a clean case of the market pricing finite supply rather than quality.

Strategic Implications

Read across categories, the findings point the same direction. Premium is paid for what cannot be reproduced, and the catalogue's job is to make that scarcity visible and believable to the room. A few practical things follow.

Most of it is decided at sourcing. The biggest premiums sit on features a specialist cannot create after the fact: an exotic dial, a steel case, a closed distillery, a Himalaya skin. That argues for putting consignment effort behind the rare configurations rather than chasing volume, since one well-sourced Paul Newman Daytona or salmon Calatrava does more for a sale than a dozen ordinary lots beside it.

What cannot be sourced can often be documented, and that is the more controllable half. A Patek extract is worth around 24%, original papers on a Daytona near 47%, a direct-from-Domaine line on a DRC lot about 22%. None of these need a rarer object, only the work to get the paperwork and present it properly.

Routing should resist a house default. No city pays more in general; handbags do better in Hong Kong, vintage sports watches in New York. Deciding venue by category rather than by policy picks up a premium a blanket rule would miss. And estimates deserve a defensive hand: condition is asymmetric, a grade lost costs more than a grade gained, so anything below top condition warrants conservative pricing to keep lots clearing.

None of this replaces a specialist's judgment. The model only sees what a catalogue records and a saleroom prices. It does not see the relationship that wins a single-owner collection, the freshness a long-held lot carries, or the timing of a market this study treats as fixed. What it offers is narrower and, within those limits, reliable: a measured account of which recorded attributes have moved hammer prices, and by how much.

Appendix

Method. Premium ratio is hammer price divided by mid-estimate, the average of published low and high. Each category model is an ordinary least squares regression of log-premium on one-hot encoded categorical attributes, with reference levels set to the most common configuration so coefficients read as premium relative to a typical lot. Using log-premium keeps a beat and a miss of equal proportion symmetric. Confidence intervals are at the 95% level; where a coefficient rests on few lots, that is noted in the text rather than hidden. Sample sizes range from roughly 150 to 320 sold lots per category.

Data. The figures in this report are illustrative, generated to demonstrate the analytical method on realistic category structures while the 2026 Luxury Week results are still pending. The schema mirrors Sotheby's Past Auctions across New York, Hong Kong, Geneva, and London, 2021 to 2026, recording each lot's sale date and location, category-specific attributes, published estimates, hammer price, and buyer region where disclosed. Macro buyer-share figures cited in Section III are real, drawn from Sotheby's December 2022 press release.

Code and full report. The complete pipeline, the interactive version of every figure, and additional category studies (DRC Burgundy and Japanese whisky) are available in the online report and project repository.

The Sotheby's Luxury Week Playbook · Independent Market Intelligence Report · June 2026